

Production Readiness Roadmap — Persistent AI Systems

Mapping every SDE-I JD requirement to concrete improvements on your project.

No new features — purely making the existing codebase industry-grade.



"What You'll Do" → What You CAN Do

1. SOLID Principles & DRY (✓ Already Started, Needs More)

You already have [renderer_base.py](#) from your SOLID refactoring conversation. What's left:

Principle	Current State	Action
S — Single Responsibility	main.py does CLI + orchestration + validation	Extract is_valid_response() into a validators.py module
O — Open/Closed	Renderer is hardcoded to Gemini	Create an abstract LLMPProvider interface so you can swap models without changing pipeline code
L — Liskov Substitution	renderer.py and renderer_streaming.py don't share a base class	Both should inherit from renderer_base.py and be interchangeable
I — Interface Segregation	semantic_search.py does embedding + indexing + searching	Split into embedder.py , indexer.py , searcher.py
D — Dependency Inversion	main.py imports concrete classes directly	Use dependency injection — pass objects in, don't hardcode imports
DRY	API call logic duplicated across renderer + summarizer	Already addressed with renderer_base.py , verify no remaining duplication

2. CI/CD Pipeline (GitHub Actions) — NEW Huge Resume Point

You're already on GitHub. Set up:

```

# .github/workflows/ci.yml
name: CI Pipeline
on: [push, pull_request]
jobs:
  test:
    runs-on: ubuntu-latest
    steps:
      - uses: actions/checkout@v4
      - uses: actions/setup-python@v5
        with:
          python-version: '3.11'
      - run: pip install -r requirements.txt
      - run: pip install pytest pytest-cov
      - run: pytest tests/ --cov=. --cov-report=xml

  lint:
    runs-on: ubuntu-latest
    steps:
      - uses: actions/checkout@v4
      - run: pip install ruff
      - run: ruff check .

```

What this demonstrates to employers:

- You understand automated quality gates
 - You know GitHub Actions (mentioned in the JD explicitly)
 - Every push is automatically validated
-

3. Observability — Logging, Metrics, Tracing

Your current logging is just `print()` statements. Industry-grade means:

What	How	Why
Structured Logging	Replace all <code>print()</code> with Python's <code>logging</code> module	Severity levels (DEBUG/INFO/WARNING/ERROR), log to files, timestamps
Metrics	Add counters: response times, API failures, cache hits, memory operations	Shows you understand observability
Health Check	Add a <code>/health</code> endpoint (now you actually use	Standard practice for any service

What	How	Why
	Flask!)	
Error Tracking	Structured error handling with error codes, not just <code>except Exception as e</code>	Specific exception types, meaningful error messages

```
# Example: What your logging should look like
import logging
logger = logging.getLogger(__name__)

# Instead of: print("    >> [Traffic Control] AI response invalid.")
logger.warning("Traffic control: discarded invalid AI response",
               extra={"response_length": len(response), "cycle": cycle_number})
```

4. Testing Framework (as discussed)

Layer	What	Coverage Target
Unit Tests	Each module independently (mock external deps)	80%+
Integration Tests	Full pipeline with mocked API	Key happy paths + error paths
Fixtures	Temp SQLite DB, mock Gemini responses, test lore files	Shared via <code>conftest.py</code>

5. Code Review Readiness

- **Type hints** on all functions (you have some, but not consistently)
- **Docstrings** on all public methods (Google style)
- `ruff` for linting + formatting (replaces flake8/black/isort in one tool)
- `pyproject.toml` to centralize project config (replace [setup.py](#))



"What We're Looking For" → **What You CAN Do**

6. Cloud Fluency — Containerize with Docker

```
# Dockerfile
FROM python:3.11-slim
WORKDIR /app
COPY requirements.txt .
RUN pip install --no-cache-dir -r requirements.txt
COPY . .
CMD ["python", "main.py"]
```

Plus a `docker-compose.yml` for local dev. This shows:

- You understand containerization (Docker — mentioned in JD)
- Your app is deployable anywhere
- Reproducible environments

7. Security Mindset (DevSecOps)

Your current code has security issues. Fix them to demonstrate security awareness:

Issue	Current	Fix
Secrets in repo	API_KEY.txt is likely committed	Use environment variables (<code>os.environ["GEMINI_API_KEY"]</code>) + <code>.env</code> file + add to .gitignore
Input validation	No sanitization on user input	Validate/sanitize before passing to API
Dependency scanning	None	Add <code>pip-audit</code> to CI pipeline (checks for known CVEs)
SQL injection	FTS5 queries may be vulnerable	Use parameterized queries (verify memory.py)
<code>.gitignore</code>	May be missing sensitive files	Ensure <code>*.db</code> , <code>*.json</code> (with secrets), <code>data/</code> are excluded

8. Microservices & API Design

Refactor [main.py](#) into a proper REST API:

```
POST /api/chat      → Send message, get response  
GET /api/health    → Service health check  
GET /api/memories → List memories (what manage_memory.py does)  
DELETE /api/memories/:id → Delete a memory
```

This turns your CLI tool into a **deployable service** and demonstrates REST API design.

🌟 Bonus Points → What You CAN Do

9. AI/ML Engineering (You Already Have This! 🎯)

Your project IS an agentic AI application. Document it properly:

- You're integrating LLM APIs (Gemini) for business logic
- You have RAG (Retrieval-Augmented Generation) with FAISS + memory
- You have an agentic loop with state management
- Highlight these in your README — this is exactly what the "Bonus Points" section asks for

10. Infrastructure as Code

- `Dockerfile` + `docker-compose.yml` (containerization)
- GitHub Actions CI/CD (automation)
- Environment variable management (`.env.example`)

11. Cost Awareness (FinOps)

You already have some of this:

- **Response caching** — avoids redundant API calls
 - **Token-saving** — proximity block disappears when unchanged
 - Document these decisions! Add a `DESIGN_DECISIONS.md` explaining WHY you cache, WHY you do 5-turn cycles (token efficiency), etc.
-

📋 Priority Order (What to Do First)

Priority	Task	Effort	Impact
🔴 P0	pytest test suite	2-3 hrs	Foundation for everything else

Priority	Task	Effort	Impact
● P0	Structured logging (replace print)	1 hr	Instant professionalism upgrade
● P0	Fix secrets management (API key)	30 min	Security 101
● P1	GitHub Actions CI/CD	1 hr	Automation showcase
● P1	Dockerfile + docker-compose	1 hr	Cloud readiness
● P1	Type hints + docstrings everywhere	1-2 hrs	Code review readiness
● P1	<code>pyproject.toml</code> + ruff linting	30 min	Modern Python packaging
● P2	REST API layer (Flask)	2-3 hrs	Microservices demo
● P2	Abstract LLM provider interface	1 hr	SOLID showcase
● P2	<code>DESIGN_DECISIONS.md</code>	1 hr	Shows cost/trade-off thinking